



NVGRAPH LIBRARY USER'S GUIDE

DU-08010-001_v10.2 | November 2019



TABLE OF CONTENTS

| | |
|---|----------|
| Chapter 1. Introduction..... | 1 |
| Chapter 2. nvGRAPH API Reference..... | 3 |
| 2.1. Return value nvgraphStatus_t..... | 3 |
| 2.2. nvGRAPH graph topology types..... | 4 |
| 2.3. nvGRAPH topology structure types..... | 4 |
| 2.3.1. nvgraphCSCTopology32I_t..... | 5 |
| 2.3.2. nvgraphCSRTopology32I_t..... | 5 |
| 2.3.3. nvgraphCOOTopology32I_t..... | 6 |
| 2.4. Function nvgraphGetProperty()..... | 6 |
| 2.5. Function nvgraphCreate()..... | 7 |
| 2.6. Function nvgraphDestroy()..... | 7 |
| 2.7. Function nvgraphCreateGraphDescr()..... | 7 |
| 2.8. Function nvgraphDestroyGraphDescr()..... | 8 |
| 2.9. Function nvgraphSetGraphStructure()..... | 8 |
| 2.10. Function nvgraphGetGraphStructure()..... | 9 |
| 2.11. Function nvgraphConvertTopology()..... | 10 |
| 2.12. Function nvgraphConvertGraph()..... | 11 |
| 2.13. Function nvgraphAllocateEdgeData()..... | 12 |
| 2.14. Function nvgraphSetEdgeData()..... | 13 |
| 2.15. Function nvgraphGetEdgeData()..... | 14 |
| 2.16. Function nvgraphAllocateVertexData()..... | 14 |
| 2.17. Function nvgraphSetVertexData()..... | 15 |
| 2.18. Function nvgraphGetVertexData()..... | 16 |
| 2.19. Function nvgraphExtractSubgraphByVertex()..... | 16 |
| 2.20. Function nvgraphExtractSubgraphByEdge()..... | 17 |
| 2.21. Function nvgraphWidestPath()..... | 18 |
| 2.22. Function nvgraphSssp()..... | 19 |
| 2.23. Function nvgraphSrSpmv()..... | 20 |
| 2.24. Function nvgraphPagerank()..... | 21 |
| 2.25. Function nvgraphTriangleCount()..... | 23 |
| 2.26. Function nvgraphStatusGetString()..... | 24 |
| 2.27. nvGRAPH Spectral Clustering API..... | 24 |
| 2.27.1. Structure SpectralClusteringParameter..... | 25 |
| 2.27.2. Function nvgraphSpectralClustering()..... | 26 |
| 2.27.3. Function nvgraphAnalyzeClustering()..... | 27 |
| 2.28. nvGRAPH Traversal API..... | 28 |
| 2.28.1. nvGRAPH traversal algorithms..... | 28 |
| 2.28.2. nvgraphTraversalParameter_t parameter struct..... | 28 |
| 2.28.3. nvgraphTraversalParameterInit()..... | 28 |
| 2.28.4. nvgraphTraversalSetDistancesIndex()..... | 29 |

| | |
|---|-----------|
| 2.28.5. nvgraphTraversalGetDistancesIndex()..... | 29 |
| 2.28.6. nvgraphTraversalSetPredecessorsIndex()..... | 30 |
| 2.28.7. nvgraphTraversalGetPredecessorsIndex()..... | 30 |
| 2.28.8. nvgraphTraversalSetEdgeMaskIndex()..... | 31 |
| 2.28.9. nvgraphTraversalGetEdgeMaskIndex()..... | 31 |
| 2.28.10. nvgraphTraversalSetUndirectedFlag()..... | 31 |
| 2.28.11. nvgraphTraversalGetUndirectedFlag()..... | 32 |
| 2.28.12. nvgraphTraversalSetAlpha()..... | 32 |
| 2.28.13. nvgraphTraversalGetAlpha()..... | 33 |
| 2.28.14. nvgraphTraversalSetBeta()..... | 33 |
| 2.28.15. nvgraphTraversalGetBeta()..... | 34 |
| 2.28.16. nvgraphTraversal()..... | 34 |
| Chapter 3. nvGRAPH Code Examples..... | 36 |
| 3.1. nvGRAPH convert topology example..... | 37 |
| 3.2. nvGRAPH convert graph example..... | 38 |
| 3.3. nvGRAPH pagerank example..... | 39 |
| 3.4. nvGRAPH SSSP example..... | 40 |
| 3.5. nvGRAPH Semi-Ring SPMV example..... | 41 |
| 3.6. nvGRAPH Triangles Counting example..... | 42 |
| 3.7. nvGRAPH Traversal example..... | 43 |

Chapter 1.

INTRODUCTION

Data analytics is a growing application of high-performance computing. Many advanced data analytics problems can be couched as graph problems. In turn, many of the common **graph problems** today can be couched as **sparse linear algebra**. This is the motivation for nvGRAPH, new in NVIDIA® CUDA™ 8.0, which harnesses the power of GPUs for linear algebra to handle the largest graph analytics and big data analytics problems.

This release provides graph **construction** and **manipulation** primitives, and a set of useful graph **algorithms** optimized for the GPU. The core functionality is a SPMV (sparse matrix vector product) using a semi-ring model with automatic load balancing for any sparsity pattern. For more information on semi-rings and their uses, we recommend the book "**Graph Algorithms in the Language of Linear Algebra**", by Jeremy Kepner and John Gilbert.

To use nvGRAPH you should be sure the nvGRAPH library is in the environment (**PATH on Windows, LD_LIBRARY_PATH on Linux**), **"#include nvgraph.h"** to your source files referencing the nvGRAPH API, and link your code using **-lnvgraph** on the command line, or add **libnvgraph** to your library dependencies. We have tested nvGRAPH using GCC 4.8 and higher on Linux, Visual Studio 2012 and Visual Studio 2014 on Windows.

A typical workflow for using nvGRAPH is to begin by calling **nvgraphCreate()** to initialize the library. Next the user can proceed to **upload graph data** to the library through nvGRAPH's API; if there is already a graph loaded in device memory then you just need a pointer to the data arrays for the graph. Graphs may be uploaded using the **CSR (compressed sparse row) format** and the **CSC (compressed column storage) format**, using **nvgraphCreateGraphDescr()**. This creates an opaque handle to the graph object, called the "graph descriptor", which represents the graph topology and its data. Graph data can be attached to vertices and/or edges of the graph, using **nvgraphSetVertexData()** and **nvgraphSetEdgeData()** respectively. Multiple values for data can co-exist on each edge or vertex at the same time, each is accessed by an index into the array of data sets. Then the **user can execute graph algorithms on the data, extract a subgraph from the data, or reformat the data** using the nvGRAPH API. The user can download the results back to host, or copy them to another location on the device, and once all calculations are done the user should call **nvgraphDestroy()** to free resources used by nvGRAPH.

nvGRAPH depends on features only present in CUDA capability 3.0 and higher architectures. This means that nvGRAPH will only run on Kepler generation or newer cards. This choice was made to provide the best performance possible.

We recommend the user start by inspecting the example codes provided, and adapt from there for their own use.

Chapter 2.

NVGRAPH API REFERENCE

This chapter specifies the behavior of the nvGRAPH library functions by describing their input/output parameters, data types, and error codes.

2.1. Return value `nvgraphStatus_t`

All nvGRAPH Library return values except for **NVGRAPH_STATUS_SUCCESS** indicate that the current API call failed and the user should reconfigure to correct the problem. The possible return values are defined as follows:

Return Values

| | |
|--|--|
| NVGRAPH_STATUS_SUCCESS | nvGRAPH operation was successful |
| NVGRAPH_STATUS_NOT_INITIALIZED | <p>The nvGRAPH library was not initialized. This is usually caused by the lack of a prior call, an error in the CUDA Runtime API called by the nvGRAPH routine, or an error in the hardware setup.</p> <p>To correct: call <code>nvgraphCreate()</code> prior to the function call; and check that the hardware, an appropriate version of the driver, and the nvGRAPH library are correctly installed.</p> |
| NVGRAPH_STATUS_ALLOC_FAILED | Resource allocation failed inside the nvGRAPH library. This is usually caused by a <code>cudaMalloc()</code> failure. |
| NVGRAPH_STATUS_INVALID_VALUE | <p>An unsupported value or parameter was passed to the function</p> <p>To correct: ensure that all the parameters being passed have valid values.</p> |
| NVGRAPH_STATUS_ARCH_MISMATCH | <p>The function requires a feature absent from the device architecture.</p> <p>To correct: compile and run the application on a device with appropriate compute capability.</p> |
| NVGRAPH_STATUS_MAPPING_ERROR | An access to GPU memory space failed. |
| NVGRAPH_STATUS_EXECUTION_FAILED | The GPU program failed to execute. This is often caused by a launch failure of the kernel on the GPU, which can be caused by multiple reasons. |

| | |
|--|--|
| | To correct: check that the hardware, an appropriate version of the driver, and the nvGRAPH library are correctly installed. |
| <code>NVGRAPH_STATUS_INTERNAL_ERROR</code> | An internal nvGRAPH operation failed. To correct: check that the hardware, an appropriate version of the driver, and the nvGRAPH library are correctly installed. Also, check that the memory passed as a parameter to the routine is not being deallocated prior to the routine's completion. |
| <code>NVGRAPH_STATUS_TYPE_NOT_SUPPORTED</code> | The type is not supported by this function. This is usually caused by passing an invalid graph descriptor to the function. |
| <code>NVGRAPH_STATUS_NOT_CONVERGED</code> | An algorithm failed to converge. To correct: ensure that all the parameters being passed have valid values for this algorithm, increase the maximum number of iteration and/or the tolerance. |

2.2. nvGRAPH graph topology types

nvGRAPH separates the topology (connectivity) of a graph from the values. To make specifying a topology easier, nvGRAPH supports three topology types. Each topology type defines its own storage format, which have benefits for some operations but detriments for others. Some algorithms can work only with specific topology types, see the algorithms descriptions for the list of supported topologies.

```
typedef enum
{
    NVGRAPH_CSR_32 = 0,
    NVGRAPH_CSC_32 = 1,
    NVGRAPH_COO_32 = 2
} nvgraphTopologyType_t;
```

Topology types

| | |
|-----------------------------|--|
| <code>NVGRAPH_CSR_32</code> | Compressed Sparse Row format (row major format). Used in SrSPMV algorithm. Use <code>nvgraphCSRTopology32I_t</code> topology structure for this format. |
| <code>NVGRAPH_CSC_32</code> | Compressed Sparse Column format (column major format). Used in SSSP , WidestPath and Pagerank algorithms. Use <code>nvgraphCSCTopology32I_t</code> topology structure for this format. |
| <code>NVGRAPH_COO_32</code> | Coordinate list format with source or destination major. Not used in any algorithm and provided for data storage only. Use <code>nvgraphCOOTopology32I_t</code> topology structure for this format. |

2.3. nvGRAPH topology structure types

Graph topology structures are used to set or retrieve topology data. Users should use the structure that corresponds to the chosen topology type.

2.3.1. nvgraphCSCTopology32I_t

Used for **NVGRAPH_CSC_32** topology type

```
struct nvgraphCSCTopology32I_st {
    int nvertices;
    int nedges;
    int *destination_offsets;
    int *source_indices;
};
typedef struct nvgraphCSCTopology32I_st *nvgraphCSCTopology32I_t;
```

Structure fields

| | |
|----------------------------|---|
| nvertices | Number of vertices in the graph. |
| nedges | Number of edges in the graph. |
| destination_offsets | Array of size nvertices +1, where i element equals to the number of the first edge for this vertex in the list of all incoming edges in the source_indices array. Last element stores total number of edges |
| source_indices | Array of size nedges , where each value designates source vertex for an edge. |

2.3.2. nvgraphCSRTopology32I_t

Used for **NVGRAPH_CSR_32** topology type

```
struct nvgraphCSRTopology32I_st {
    int nvertices;
    int nedges;
    int *source_offsets;
    int *destination_indices;
};
typedef struct nvgraphCSRTopology32I_st *nvgraphCSRTopology32I_t;
```

Structure fields

| | |
|----------------------------|--|
| nvertices | Number of vertices in the graph. |
| nedges | Number of edges in the graph. |
| source_offsets | Array of size nvertices +1, where i element equals to the number of the first edge for this vertex in the list of all outgoing edges in the destination_indices array. Last element stores total number of edges |
| destination_indices | Array of size nedges , where each value designates destination vertex for an edge. |

2.3.3. nvgraphCOOTopology32I_t

Used for **NVGRAPH_COO_32** topology type

```
struct nvgraphCOOTopology32I_st {
    int nvertices;
    int nedges;
    int *source_indices;
    int *destination_indices;
    nvgraphTag_t tag;
};
typedef struct nvgraphCOOTopology32I_st *nvgraphCOOTopology32I_t;
```

Structure fields

| | |
|----------------------------|--|
| nvertices | Number of vertices in the graph. |
| nedges | Number of edges in the graph. |
| source_indices | Array of size nedges , where each value designates the source vertex for an edge. |
| destination_indices | Array of size nedges , where each value designates the destination vertex for an edge. |
| tag | One of the values of NVGRAPH_UNSORTED , NVGRAPH_SORT_BY_SOURCE or NVGRAPH_SORT_BY_DESTINATION to indicate topology order. |

2.4. Function nvgraphGetProperty()

```
nvgraphStatus_t
nvgraphGetProperty(libraryPropertyType type, int *value);
```

Returns property value of a library, such as version number.

Input

| | |
|-------------|---|
| type | Identifier of a library property, such as MAJOR_VERSION , MINOR_VERSION or PATCH_LEVEL |
|-------------|---|

Output

| | |
|--------------|-------------------------------|
| value | Value of a requested property |
|--------------|-------------------------------|

Return Values

| | |
|-------------------------------------|--|
| NVGRAPH_STATUS_SUCCESS | nvGRAPH successfully created the handle. |
| NVGRAPH_STATUS_INVALID_VALUE | Invalid property requested. |

2.5. Function `nvgraphCreate()`

```
nvgraphStatus_t
nvgraphCreate(nvgraphHandle_t *handle);
```

Creates only an opaque handle, and allocates small data structures on the host. This handle is used in all of the nvGRAPH functions, so this function should be called first, before any other calls are made to the library.

Input/Output

| | |
|---------------|--|
| handle | Pointer to a <code>nvgraphHandle_t</code> object |
|---------------|--|

Return Values

| | |
|--|--|
| <code>NVGRAPH_STATUS_SUCCESS</code> | nvGRAPH successfully created the handle. |
| <code>NVGRAPH_STATUS_ALLOC_FAILED</code> | The allocation of resources for the handle failed. |
| <code>NVGRAPH_STATUS_INTERNAL_ERROR</code> | An internal driver error was detected. |

2.6. Function `nvgraphDestroy()`

```
nvgraphStatus_t
nvgraphDestroy(nvgraphHandle_t handle);
```

Destroys a handle created with `nvgraphCreate()`. This will automatically release any allocated memory objects created with this handle, for example any graphs and their vertices' and edges' data. Any subsequent usage of this handle after calling `nvgraphDestroy()` will be invalid. Any calls to the nvGRAPH API after `nvgraphDestroy()` is called will return 'NVGRAPH_UNINITIALIZED' errors.

Input

| | |
|---------------|--|
| handle | The <code>nvgraphHandle_t</code> object of the handle to be destroyed. |
|---------------|--|

Return Values

| | |
|---|--|
| <code>NVGRAPH_STATUS_SUCCESS</code> | nvGRAPH successfully destroyed the handle. |
| <code>NVGRAPH_STATUS_INVALID_VALUE</code> | The <code>handle</code> parameter is not a valid handle. |

2.7. Function `nvgraphCreateGraphDescr()`

```
nvgraphStatus_t
nvgraphCreateGraphDescr(nvgraphHandle_t handle, nvgraphGraphDescr_t *descrG);
```

Creates opaque handle for a graph structure. This handle is required for any operation on the graph.

Input

| | |
|---------------|------------------------|
| handle | nvGRAPH library handle |
|---------------|------------------------|

Input/Output

| | |
|---------------|---|
| descrG | Pointer to the empty <code>nvgraphGraphDescr_t</code> structure object. |
|---------------|---|

Return Values

| | |
|-------------------------------------|-----------------------------------|
| NVGRAPH_STATUS_SUCCESS | Success |
| NVGRAPH_STATUS_INVALID_VALUE | Bad library handle is provided |
| NVGRAPH_STATUS_ALLOC_FAILED | Cannot allocate graph descriptor. |

2.8. Function `nvgraphDestroyGraphDescr()`

```
nvgraphStatus_t
nvgraphDestroyGraphDescr(nvgraphHandle_t handle, nvgraphGraphDescr_t descrG);
```

Destroys a graph handle created with `nvgraphCreateGraphDescr()`. This won't release any memory allocated for this graph until the nvGRAPH library handle is destroyed. Calls to manipulate destroyed graphs will return `NVGRAPH_STATUS_INVALID_VALUE`.

Input

| | |
|---------------|---------------------------------|
| handle | nvGRAPH library handle |
| descrG | Graph descriptor to be released |

Return Values

| | |
|--|---|
| NVGRAPH_STATUS_SUCCESS | Successful release of the graph descriptor |
| NVGRAPH_STATUS_TYPE_NOT_SUPPORTED | Graph is stored with unknown type of data |
| NVGRAPH_STATUS_INVALID_VALUE | Invalid library handle or graph descriptor handle |

2.9. Function `nvgraphSetGraphStructure()`

```
nvgraphStatus_t
nvgraphSetGraphStructure( nvgraphHandle_t handle, nvgraphGraphDescr_t descrG,
    void* topologyData, nvgraphTopologyType_t TType);
```

This call sets both topology data and topology type for the given graph descriptor. Graph topology should be set only once. Users should choose one of the supported topologies, fill in the corresponding structure for the graph structure initialization and provide a pointer to this structure. The topology data and type are given in parameters **topologyData** and **TType**. Typically graph topology data includes a number of

vertices, number of edges and connectivity information. Look at the description of the corresponding topology structures for details.

Input

| | |
|---------------------|--|
| handle | nvGRAPH library handle |
| topologyData | Pointer to a filled structure of one of the types { <code>nvgraphCSRTopology32I_t</code> , <code>nvgraphCSCTopology32I_t</code> }. The particular type to be used is defined by parameter <code>TType</code> . |
| TType | Graph topology type. This value should be equal to one of the possible values of the enum <code>nvgraphTopologyType_t</code> . This defines what data structure should be provided by the <code>topologyData</code> parameter. |

Input/Output

| | |
|---------------|--|
| descrG | Graph descriptor. Must not have topology previously defined. |
|---------------|--|

Return Values

| | |
|--|---|
| <code>NVGRAPH_STATUS_SUCCESS</code> | Success |
| <code>NVGRAPH_STATUS_INVALID_VALUE</code> | Invalid library handle, topology data structure pointer or topology values, or graph topology was already set |
| <code>NVGRAPH_STATUS_TYPE_NOT_SUPPORTED</code> | Provided topology type is not supported |
| <code>NVGRAPH_STATUS_INTERNAL_ERROR</code> | Unknown internal error was caught |

2.10. Function `nvgraphGetGraphStructure()`

```
nvgraphStatus_t
nvgraphGetGraphStructure( nvgraphHandle_t handle, nvgraphGraphDescr_t descrG,
    void* topologyData, nvgraphTopologyType_t* TType);
```

This function allows users to retrieve a given graph's topology and topology data such as the number of vertices and edges in the graph. Users must provide a graph descriptor as well as an empty topology structure, where this information will be stored.

Input

| | |
|---------------|---|
| handle | nvGRAPH library handle |
| descrG | Graph descriptor. This graph should have its topology structure previously set. |

Input/Output

| | |
|---------------------|---|
| topologyData | Pointer to a structure of one of the types { <code>nvgraphCSRTopology32I_t</code> , <code>nvgraphCSCTopology32I_t</code> } matching the graph topology. If the field is <code>NULL</code> it will be ignored and only graph topology will be returned. The user can point source and destination fields in the structure to a host or device memory to retrieve |
|---------------------|---|

| | |
|--------------|---|
| | connectivity information, if one of the fields is NULL it will be ignored. |
| TType | Pointer to <code>nvgraphTopologyType_t</code> where graph topology will be returned. If the field is NULL it will be ignored and only topology data will be returned. |

Return Values

| | |
|--|--|
| NVGRAPH_STATUS_SUCCESS | Success |
| NVGRAPH_STATUS_INVALID_VALUE | Invalid library handle, graph descriptor or topology for the graph is not set. |
| NVGRAPH_STATUS_TYPE_NOT_SUPPORTED | Unsupported topology or graph's topology doesn't match provided parameter. |

2.11. Function `nvgraphConvertTopology()`

```
nvgraphStatus_t
nvgraphConvertTopology(nvgraphHandle_t handle,
    nvgraphTopologyType_t srcTType, void *srcTopology, void *srcEdgeData,
    cudaDataType_t *dataType,
    nvgraphTopologyType_t dstTType, void *dstTopology, void *dstEdgeData);
```

Convert one of the supported topologies to another. In case the source and destination topologies are the same, the function will perform a straight memory copy.

This function assumes that source and destination arrays within the topologies and edge data reside in GPU (device) memory. It is the user's responsibility to allocate memory and copy data between GPU memory and CPU memory using standard CUDA runtime API routines, such as `cudaMalloc()`, `cudaFree()`, `cudaMemcpy()`, and `cudaMemcpyAsync()`.

If the destination topology is `nvgraphCOOTopology32I_st`, the tag field needs to be set to one of the values of `NVGRAPH_UNSORTED`, `NVGRAPH_SORT_BY_SOURCE` or `NVGRAPH_SORT_BY_DESTINATION` to tell the library how the user wants entries sorted (or not).

The function requires extra buffer storage in the device memory for some of the conversion operations, the extra storage is proportional to the topology size. The function is executed asynchronously with respect to the host and may return control to the application on the host before the result is ready.

Input

| | |
|--------------------|---|
| handle | nvGRAPH library handle. |
| srcTType | Source topology type. This value should be equal to one of the possible values of the enum <code>nvgraphTopologyType_t</code> . This defines what data structure type should be provided by the <code>srcTopology</code> parameter. |
| srcTopology | Pointer to a structure of one of the types <code>{nvgraphCSRTopology32I_t, nvgraphCSTopology32I_t, nvgraphCOOTopology32I_st}</code> . The particular type to |

| | |
|--------------------|--|
| | be used is defined by parameter srcTType . The function assumes that source and destination arrays within the structure reside in device memory. |
| srcEdgeData | Pointer to the user memoryspace where edge data are stored. Must be device memory. |
| dataType | Edge data type. This value should be equal to one of <code>CUDA_R_32F</code> or <code>CUDA_R_64F</code> . This defines what data type is provided by the srcEdgeData and dstEdgeData parameters. |
| dstTType | Destination topology type. This value should be equal to one of the possible values of the enum <code>nvgraphTopologyType_t</code> . This defines what data structure type should be provided by the dstTopology and parameter. |

Input/Output

| | |
|--------------------|--|
| dstTopology | Pointer to a structure of one of the types <code>{nvgraphCSRTopology32I_t, nvgraphCSTopology32I_t, nvgraphCOOTopology32I_st}</code> where conversion results will be stored. The particular type to be used is defined by parameter dstTType . The function assumes that source and destination arrays within the structure are pre-allocated by the user with at least the number of bytes needed by the topology. Source and destination arrays are assumed to reside in device memory. |
| dstEdgeData | Pointer to the user memoryspace where converted edge data will be stored. Must be device memory and have at least <code>number_of_vertices*sizeof(dataType)</code> bytes. |

Return Values

| | |
|--|--|
| <code>NVGRAPH_STATUS_SUCCESS</code> | Success. |
| <code>NVGRAPH_STATUS_INVALID_VALUE</code> | Bad parameter(s). |
| <code>NVGRAPH_STATUS_TYPE_NOT_SUPPORTED</code> | An internal operation failed. |
| <code>NVGRAPH_STATUS_INTERNAL_ERROR</code> | The type of at least one topology or edge set is not supported. Currently we support float and double type values. |

2.12. Function `nvgraphConvertGraph()`

```
nvgraphStatus_t
nvgraphConvertGraph(nvgraphHandle_t handle,
    nvgraphGraphDescr_t srcDescrG, nvgraphGraphDescr_t dstDescrG,
    nvgraphTopologyType_t dstTType);
```

Convert one of the supported graph types to another. The function will allocate the necessary memory for the destination graph. It is recommended to use this function over **nvgraphConvertTopology** when converting a large data set.

In addition to the destination graph memory, the function requires extra buffer storage in the device memory for some of the conversion operation, the extra storage is

proportional to the topology size. The function is executed asynchronously with respect to the host and may return control to the application on the host before the result is ready.

Input

| | |
|------------------|--|
| handle | nvGRAPH library handle. |
| srcDescrG | Source graph descriptor. This graph should have its topology structure previously set, and optionally vertex and edge data set. |
| dstTType | Destination topology type. This value should be equal to one of the possible values of { <code>nvgraphCSRTopology32I_t</code> , <code>nvgraphCSCTopology32I_t</code> , <code>nvgraphCOOTopology32I_t</code> }. This defines what data structure type should be provided by the <code>dstTopology</code> and parameter. |

Input/Output

| | |
|------------------|--|
| dstDescrG | Destination graph descriptor. Must be an empty descriptor created with <code>nvgraphCreateGraphDescr</code> and not have a topology set. |
|------------------|--|

Return Values

| | |
|--|--|
| NVGRAPH_STATUS_SUCCESS | Success. |
| NVGRAPH_STATUS_INVALID_VALUE | Bad parameter(s). |
| NVGRAPH_STATUS_TYPE_NOT_SUPPORTED | An internal operation failed. |
| NVGRAPH_STATUS_INTERNAL_ERROR | The type of at least one topology or edge set is not supported. Currently we support float and double type values. |

2.13. Function `nvgraphAllocateEdgeData()`

```
nvgraphStatus_t
nvgraphAllocateEdgeData(nvgraphHandle_t handle, nvgraphGraphDescr_t descrG,
    size_t numsets, cudaDataType_t *settypes);
```

Allocates one or more storages for the data associated with graph edges. Number of allocated storages is specified by the `numsets` parameter. Types for each of the allocated storages should be provided in the array `settypes` of size `numsets`. Right now nvGRAPH graphs are limited to have data storages to have same type and same size - all elements of `settypes` array should be the same and all of those storages will have number of elements equal to the number of edges in the graph. Vertices data allocated with `nvgraphAllocateVerticesData()` function should have the same datatype as edge data. These storages could later be used in other functions by indices from **0** to **numsets-1**. This function could be called successfully only once for each graph.

Input

| | |
|---------------|------------------------|
| handle | nvGRAPH library handle |
|---------------|------------------------|

| | |
|-----------------|---|
| numsets | Number of datasets to allocate for the edges. Should be more than 0. |
| settypes | Array of size numsets that specifies types of allocated datasets. All values in this array should be the same and match graph's datasets data type, if exists. |

Input/Output

| | |
|---------------|---|
| descrG | Descriptor of the graph for which edge data is allocated. Should not have previously allocated edge data and have it's topology properly initialized. |
|---------------|---|

Return Values

| | |
|--|--|
| NVGRAPH_STATUS_SUCCESS | Success. |
| NVGRAPH_STATUS_INVALID_VALUE | Invalid function parameters, inconsistent types in the type array, types doesn't match graph's type or graph is not initialized for data allocation. |
| NVGRAPH_STATUS_TYPE_NOT_SUPPORTED | Types provided in parameter are not supported. |

2.14. Function `nvgraphSetEdgeData()`

```
nvgraphStatus_t
nvgraphSetEdgeData(nvgraphHandle_t handle, nvgraphGraphDescr_t descrG,
void *edgeData, size_t setnum);
```

Update a specific edge value set (weights) of the graph with the user's provided values.

Input

| | |
|------------------|---|
| handle | nvGRAPH library handle. |
| descrG | nvGRAPH graph descriptor, should contain the connectivity information and the edge set setnum |
| *edgeData | Pointer to the data to load into the edge value set. This entry expects to read one value for each edge. Conversions are not supported so the user's type before the void* cast should be equivalent to the one specified in nvgraphAllocateEdgeData |
| setnum | The identifier of the set to update. This assumes that setnum is one of the the edge set allocated in the past using nvgraphAllocateEdgeData . Sets have 0-based index |

Return Values

| | |
|--------------------------------------|-------------------------------|
| NVGRAPH_STATUS_SUCCESS | Success. |
| NVGRAPH_STATUS_INVALID_VALUE | Bad parameter(s). |
| NVGRAPH_STATUS_INTERNAL_ERROR | An internal operation failed. |

2.15. Function nvgraphGetEdgeData()

```
nvgraphStatus_t
nvgraphGetEdgeData(nvgraphHandle_t handle, nvgraphGraphDescr_t descrG,
void *edgeData, size_t setnum);
```

Downloads one dataset associated with graph edges using **setnum** index to the user memoryspace. **edgeData** could point to either host or device memoryspace. Size of the data transfer depends on the edges number of the graph and graph's data type.

Input

| | |
|---------------|---|
| handle | nvGRAPH library handle. |
| descrG | Graph descriptor. Graph should contain at least one data set associated with it's vertices |
| setnum | Index of the source data set of the graph edge data. Value should be between 0 and edge_dataset_number-1 |

Output

| | |
|-----------------|---|
| edgeData | Pointer to the user memoryspace where edge data will be stored. Could be either host or device memory and have at least number_of_edges*sizeof(graph_data_type) bytes. |
|-----------------|---|

Return Values

| | |
|--|--|
| NVGRAPH_STATUS_SUCCESS | Success. |
| NVGRAPH_STATUS_INVALID_VALUE | Incorrect function parameter, graph has no associated edge data sets or topology type doesn't match. |
| NVGRAPH_STATUS_TYPE_NOT_SUPPORTED | Graph datatype or topology type is not supported. |
| NVGRAPH_STATUS_INTERNAL_ERROR | Memory copy failed. |

2.16. Function nvgraphAllocateVertexData()

```
nvgraphStatus_t
nvgraphAllocateVertexData(nvgraphHandle_t handle, nvgraphGraphDescr_t descrG,
size_t numsets, cudaDataType_t *settypes);
```

Allocates one or more storages for the data associated with graph vertices. Number of allocated storages is specified by the **numsets** parameter. Types for each of the allocated storages should be provided in the array **settypes** of size **numsets**. Right now nvGRAPH graphs are limited to have data storages to have same type and same size - all elements of **settypes** array should be the same and all of those storages will have number of elements equal to the number of vertices in the graph. Edge data allocated with **nvgraphAllocateEdgeData()** function should have the same datatype as vertex data. These storages could later be used in other functions by indices from **0** to **numsets-1**. This function could be called successfully only once for each graph.

Input

| | |
|-----------------|---|
| handle | nvGRAPH library handle |
| numsets | Number of datasets to allocate for the vertices. Should be more than 0. |
| settypes | Array of size numsets that specifies types of allocated datasets. All values in this array should be the same and match graph's datasets data type, if exists. |

Input/Output

| | |
|---------------|---|
| descrG | Descriptor of the graph for which edge data is allocated. Should not have previously allocated edge data and have it's topology properly initialized. |
|---------------|---|

Return Values

| | |
|--|--|
| NVGRAPH_STATUS_SUCCESS | Success. |
| NVGRAPH_STATUS_INVALID_VALUE | Invalid function parameters, inconsistent types in the type array, types doesn't match graph's type or graph is not initialized for data allocation. |
| NVGRAPH_STATUS_TYPE_NOT_SUPPORTED | Types provided in parameter are not supported. |

2.17. Function `nvgraphSetVertexData()`

```
nvgraphStatus_t
nvgraphSetVertexData(nvgraphHandle_t handle, nvgraphGraphDescr_t descrG,
    void *vertexData, size_t setnum);
```

Update a specific vertex value set of the graph with the user's provided values.

Input

| | |
|--------------------|---|
| handle | nvGRAPH library handle. |
| descrG | nvGRAPH graph descriptor, should contain the connectivity information and vertex set setnum . |
| *vertexData | Pointer to the data to load into the vertex value set. This entry expects to read one value for each vertex. Conversions are not supported so the user's type before the void* cast should be equivalent to the one specified in nvgraphAllocateVertexData |
| setnum | The identifier of the set to update. This assumes that setnum is one of the the vertex set allocated in the past using nvgraphAllocateVertexData . Sets have 0-based index |

Return Values

| | |
|--------------------------------------|-------------------------------|
| NVGRAPH_STATUS_SUCCESS | Success. |
| NVGRAPH_STATUS_INVALID_VALUE | Bad parameter(s). |
| NVGRAPH_STATUS_INTERNAL_ERROR | An internal operation failed. |

2.18. Function nvgraphGetVertexData()

```
nvgraphStatus_t
nvgraphGetVertexData(nvgraphHandle_t handle, nvgraphGraphDescr_t descrG,
void *vertexData, size_t setnum);
```

Downloads one dataset associated with graph vertices using **setnum** index to the user memoryspace. **vertexData** could point to either host or device memoryspace. Size of the data transfer depends on the vertex number of the graph and graph's data type.

Input

| | |
|---------------|---|
| handle | nvGRAPH library handle. |
| descrG | Graph descriptor. Graph should contain at least one data set associated with it's vertices |
| setnum | Index of the source data set of the graph vertex data. Value should be between 0 and vertex_dataset_number-1 |

Output

| | |
|-----------------|--|
| edgeData | Pointer to the user memoryspace where vertex data will be stored. Could be either host or device memory and have at least number_of_vertices*sizeof(graph_data_type) bytes. |
|-----------------|--|

Return Values

| | |
|--|--|
| NVGRAPH_STATUS_SUCCESS | Success. |
| NVGRAPH_STATUS_INVALID_VALUE | Incorrect function parameter, graph has no associated vertex data sets or topology type doesn't match. |
| NVGRAPH_STATUS_TYPE_NOT_SUPPORTED | Graph datatype or topology type is not supported. |
| NVGRAPH_STATUS_INTERNAL_ERROR | Memory copy failed. |

2.19. Function nvgraphExtractSubgraphByVertex()

```
nvgraphStatus_t
nvgraphExtractSubgraphByVertex(nvgraphHandle_t handle,
nvgraphGraphDescr_t descrG, nvgraphGraphDescr_t subdescrG,
int *subvertices, size_t numvertices);
```

Create a new graph by extracting a subgraph given an array of vertices, consisting of row indices in the graph incidence matrix; array must be: (i) free of duplicates; (ii) sorted in ascending order; (iii) must consist of indices with values between 0 and **graph_nvertices-1**.

Input

| | |
|---------------|---|
| handle | nvGRAPH handle of the source graph (original graph) |
|---------------|---|

| | |
|--------------------|--|
| descrG | nvGRAPH descriptor of the source graph (original graph) |
| subvertices | array containing vertex indices (row indices in graph incidence matrix) of the subgraph to be extracted; array must be: (i) free of duplicates; (ii) sorted in ascending order; (iii) must consist of indices with values between 0 and graph_nvertices-1 . |
| numvertices | the size of subvertices[] array. Should be more than 0 and less or equal to the number of graph's vertices. |

Output

| | |
|------------------|---|
| subdescrG | nvGRAPH graph descriptor of the target graph (subgraph) |
|------------------|---|

Return Values

| | |
|--|---|
| NVGRAPH_STATUS_SUCCESS | nvGRAPH target (subgraph) was created successfully. |
| NVGRAPH_STATUS_INVALID_VALUE | Bad parameter(s). |
| NVGRAPH_STATUS_TYPE_NOT_SUPPORTED | The type of specified nvGRAPH is not supported. |

2.20. Function `nvgraphExtractSubgraphByEdge()`

```
nvgraphStatus_t
nvgraphExtractSubgraphByEdge(nvgraphHandle_t handle,
                             nvgraphGraphDescr_t descrG, nvgraphGraphDescr_t subdescrG,
                             int *subedges, size_t numedges);
```

Create a new graph by extracting a subgraph given an array of edges, consisting of indices in the `col_ind[]` array of the the graph incidence matrix CSR representation); the array of edges must be: (i) free of duplicates; (ii) sorted in ascending order; (iii) must consist of indices with values between 0 and **graph_nedges-1**.

Input

| | |
|-----------------|--|
| handle | nvGRAPH handle of the source graph (original graph) |
| descrG | nvGRAPH descriptor of the source graph (original graph) |
| subedges | array containing edge indices (indices in the <code>col_ind[]</code> array of the the graph incidence matrix CSR representation) of the subgraph to be extracted; array must be: (i) free of duplicates; (ii) sorted in ascending order; (iii) must consist of indices with values between 0 and graph_nedges-1 |
| numedges | the size of subedges[] array, should be more than 0 and less or equal to number of graph's edges |

Output

| | |
|------------------|---|
| subdescrG | nvGRAPH graph descriptor of the target graph (subgraph) |
|------------------|---|

Return Values

| | |
|-------------------------------|---|
| NVGRAPH_STATUS_SUCCESS | nvGRAPH target (subgraph) was created successfully. |
|-------------------------------|---|

| | |
|--|---|
| <code>NVGRAPH_STATUS_INVALID_VALUE</code> | Bad parameter(s). |
| <code>NVGRAPH_STATUS_TYPE_NOT_SUPPORTED</code> | The type of specified nvGRAPH is not supported. |

2.21. Function `nvgraphWidestPath()`

```
nvgraphStatus_t
nvgraphWidestPath(nvgraphHandle_t handle, const nvgraphGraphDescr_t descrG,
                  const size_t weight_index, const int *source_vert,
                  const size_t widest_path_index);
```

Find the widest path from the vertex at `source_index` to every other vertex; this problem is also known as 'the bottleneck path problem' or 'the maximum capacity path problem'.

If some vertices are unreachable, the widest path to those vertices is $-\infty$. In limited-precision arithmetic, this corresponds to `-FLT_MAX` or `-DBL_MAX` depending on the value type of the set (`CUDA_R_32F` or `CUDA_R_64F` respectively).

Input

| | |
|---------------------------|--|
| <code>handle</code> | nvGRAPH library handle. |
| <code>descrG</code> | nvGRAPH graph descriptor, should contain the connectivity information in <code>NVGRAPH_CSC_32</code> , at least 1 edge set (the capacity) and 1 vertex set (to store the result). |
| <code>weight_index</code> | Index of the edge set for the weights. |
| <code>*source_vert</code> | Index of the source, using 0-based indexes. |

Output

| | |
|--------------------------------|---|
| <code>widest_path_index</code> | <p>The values stored inside the vertex set at <code>widest_path_index</code> (<code>VertexData[widest_path_index]</code>) are the widest path values. <code>VertexData[widest_path_index][i]</code> is the length of the widest path between <code>source_vert</code> and vertex <code>i</code>. If vertex <code>i</code> is not reachable from <code>source_vert</code>, <code>VertexData[widest_path_index][i] = $-\infty$</code>.</p> <p>Users can get a copy of the result using <code>nvgraphGetVertexData</code></p> |
|--------------------------------|---|

Return Values

| | |
|--|---|
| <code>NVGRAPH_STATUS_SUCCESS</code> | Success. |
| <code>NVGRAPH_STATUS_INVALID_VALUE</code> | Bad parameter(s). |
| <code>NVGRAPH_STATUS_TYPE_NOT_SUPPORTED</code> | The type of at least one vertex or edge set is not supported. |
| <code>NVGRAPH_STATUS_INTERNAL_ERROR</code> | An internal operation failed. |

2.22. Function nvgraphSssp()

```
nvgraphStatus_t
nvgraphSssp(nvgraphHandle_t handle, const nvgraphGraphDescr_t descrG,
            const size_t weight_index, const int *source_vert,
            const size_t sssp_index);
```

The Single Source Shortest Path (SSSP) algorithm calculates the shortest path distance from a single vertex in the graph to all other vertices.

If some vertices are unreachable, the shortest path to those vertices is ∞ . In limited-precision arithmetic, that corresponds to FLT_MAX or DBL_MAX depending on the value type of the set (CUDA_R_32F or CUDA_R_64F respectively).

Input

| | |
|---------------------|--|
| handle | nvGRAPH library handle. |
| descrG | nvGRAPH graph descriptor, should contain the connectivity information in NVGRAPH_CSC_32 , at least 1 edge set (distances) and 1 vertex set (the shortest path lengths). |
| weight_index | Index of the edge set for the weights. The default value is 0, meaning the first edge set. |
| *source_vert | Index of the source, using 0-based indexes. |

Output

| | |
|-------------------|---|
| sssp_index | <p>The values stored inside the vertex set at sssp_index (VertexData[sssp_index]) are the shortest path values. VertexData[sssp_index][i] is the length of the shortest path between source_vert and vertex i. If vertex i is not reachable from source_vert, VertexData[sssp_index][i] = ∞.</p> <p>User can get a copy of the result using nvgraphGetVertexData</p> |
|-------------------|---|

Return Values

| | |
|--|---|
| NVGRAPH_STATUS_SUCCESS | Success. |
| NVGRAPH_STATUS_INVALID_VALUE | Bad parameter(s). |
| NVGRAPH_STATUS_TYPE_NOT_SUPPORTED | The type of at least one vertex or edge set is not supported. |
| NVGRAPH_STATUS_INTERNAL_ERROR | An internal operation failed. |

2.23. Function nvgraphSrSpmv()

```
nvgraphStatus_t
nvgraphSrSpmv(nvgraphHandle_t handle, const nvgraphGraphDescr_t descrG,
               const size_t weight_index, const void *alpha, const size_t x_index,
               const void *beta, const size_t y_index, const nvgraphSemiring_t SR);
```

The Semi-Ring Sparse Matrix Vector multiplication is an operation of the type $y = \alpha * A * x + \beta y$. Where :

- A is a weighted graph seen as a compressed sparse matrix in CSR, x and y are vectors, α and β are scalars
- $(*,+)$ is a set of two binary operators operating on real values and satisfying semi-ring properties.

In nvGRAPH all semi-rings operate on a set (R) with two binary operators: + and * that satisfies:

- (R, +) is associative, commutative with additive identity (additive_identity + a = a)
- (R, *) is associative with multiplicative identity (multiplicative_identity * a = a)
- Left and Right multiplication is distributive over addition
- Additive identity = multiplicative null operator (null_operator * a = a * null_operator = null_operator).

nvGRAPH's approach for sparse matrix vector multiplication on the GPU is based on the CSR MV merge-path algorithm from Duane Merrill. It is designed to handle arbitrary sparsity patterns in an efficient way by offering a good workload balance. As a result, this operation delivers consistent good performance even for networks with a power-law distribution of connections.

nvGRAPH has pre-defined useful semi-ring for graphs in nvgraphSemiring_t, so the user can select them directly.

Semi rings

| Semiring | Set | Plus | Times | Add_ident | Mult_ident |
|-----------------------|-------------------------------|------|-------|-----------|------------|
| NVGRAPH_PLUS_TIMES_SR | R | + | * | 0 | 1 |
| NVGRAPH_MIN_PLUS_SR | $R \cup \{-\infty, +\infty\}$ | min | + | ∞ | 0 |
| NVGRAPH_MAX_MIN_SR | $R \cup \{-\infty, +\infty\}$ | max | min | $-\infty$ | $+\infty$ |
| NVGRAPH_OR_AND_SR | {0,0, 1,0} | OR | AND | 0 | 1 |

Input

| | |
|---------------------|--|
| handle | nvGRAPH library handle. |
| descrG | nvGRAPH graph descriptor, should contain the connectivity information in NVGRAPH_CSR_32 , at least 1 edge set (weights) and 2 vertex sets (input vector and output vector). |
| weight_index | Index of the edge set for the weights. |
| *alpha | Scalar used for multiplication |
| x_index | Index of the vertex set for used for multiplication |
| *beta | Scalar used for multiplication. If beta is zero, the vertex set at y_index does not have to be a valid input. |
| y_index | (optional) Index of the vertex set for used for the addition. |
| SR | The semi-ring type nvgraphSemiring_t which can be NVGRAPH_PLUS_TIMES_SR , NVGRAPH_MIN_PLUS_SR , NVGRAPH_MAX_MIN_SR , NVGRAPH_OR_AND_SR . |

Output

| | |
|--------------------------|--|
| Values at y_index | <p>The values stored inside the set at y_index (VertexData[y_index]) are the result of the operation.</p> <p>User can get a copy of the result using nvgraphGetVertexData</p> |
|--------------------------|--|

Return Values

| | |
|--|---|
| NVGRAPH_STATUS_SUCCESS | Success. |
| NVGRAPH_STATUS_INVALID_VALUE | Bad parameter(s). |
| NVGRAPH_STATUS_TYPE_NOT_SUPPORTED | The type of at least one vertex or edge set is not supported. |
| NVGRAPH_STATUS_INTERNAL_ERROR | An internal operation failed. |

2.24. Function **nvgraphPagerank()**

```
nvgraphStatus_t
nvgraphPagerank(nvgraphHandle_t handle, const nvgraphGraphDescr_t descrG,
                const size_t weight_index, const void *alpha,
                const size_t bookmark_index,
                const int has_guess, const size_t pagerank_index,
                const float tolerance, const int max_iter);
```

Find the PageRank vertex values for a graph with a given transition matrix (Markov chain), a bookmark vector of dangling vertices, and the damping factor. The transition matrix is sub-stochastic (ie. each row sums to 0 or 1) and has to be provided in column major order (ie. in CSC, which is equivalent to the transposed of the sub-stochastic matrix in CSR). The bookmark vector flags vertices without outgoing edges (also called dangling vertices).

This is equivalent to an eigenvalue problem where we compute the dominant eigenpair. By construction, the maximum eigenvalue is 1, only the eigenvector is interesting. nvGRAPH computes an approximation of the Pagerank eigenvector using the power method. The number of iterations depends on the properties of the network itself; it increases when the tolerance decreases and/or alpha increases toward the limiting value of 1.

The user is free to use default values or to provide inputs for the initial guess, tolerance and maximum number of iterations.

Input

| | |
|-----------------------|---|
| handle | nvGRAPH library handle. |
| descrG | nvGRAPH graph descriptor, should contain the connectivity information in NVGRAPH_CSC_32 , at least 1 edge set and 2 vertex sets. |
| weight_index | Index of the edge set for the transition probability. |
| *alpha | The damping factor alpha represents the probability to follow an outgoing edge, standard value is 0.85. Thus 1.0-alpha is the probability to “teleport” to a random node. alpha should be greater than 0.0 and strictly lower than 1.0. |
| bookmark_index | Index of the vertex set for the bookmark of dangling nodes (VertexData[bookmark_index][i] = 1.0 if i is a dangling node, 0.0 otherwise). |
| has_guess | This parameter is used to notify nvGRAPH if it should use a user-provided initial guess. 0 means the user doesn't have a guess, in this case nvGRAPH will use a uniform vector set to 1/v . If the value is 1 nvGRAPH will read VertexData[pagerank_index] and use this as initial guess. The initial guess must not be the vector of 0s. Any value other than 1 or 0 is treated as an invalid value. |
| pagerank_index | (optional) Index of the vertex set for the initial guess if has_guess=1 |
| tolerance | Set the tolerance the approximation, this parameter should be a small magnitude value. The lower the tolerance the better the approximation. If this value is 0.0, nvGRAPH will use the default value which is 1.0E-6 . Setting too small a tolerance (less than 1.0E-6 typically) can lead to non-convergence due to numerical roundoff. Usually 0.01 and 0.0001 are acceptable. |
| max_iter | The maximum number of iterations before an answer is returned. This can be used to limit the execution time and do an early exit before the solver reaches the convergence tolerance. If this value is lower or equal to 0 nvGRAPH will use the default value, which is 500. |

Output

| | |
|---------------------------------|---|
| Values at pagerank_index | The values stored inside the vertex set at pagerank_index (VertexData[pagerank_index]) are the PageRank values. VertexData[pagerank_index][i] is the PageRank of vertex i . |
|---------------------------------|---|

| | |
|--|--|
| | Users can get a copy of the result using <code>nvgraphGetVertexData</code> |
|--|--|

Return Values

| | |
|--|--|
| <code>NVGRAPH_STATUS_SUCCESS</code> | The Pagerank iteration reached the desired <code>tolerance</code> in less than <code>max_iter</code> iterations |
| <code>NVGRAPH_STATUS_NOT_CONVERGED</code> | The Pagerank iteration did not reach the desired <code>tolerance</code> after <code>max_iter</code> iterations |
| <code>NVGRAPH_STATUS_INVALID_VALUE</code> | Bad parameter(s). |
| <code>NVGRAPH_STATUS_TYPE_NOT_SUPPORTED</code> | The type of at least one vertex or edge set is not supported. Currently we support float and double type values. |

2.25. Function `nvgraphTriangleCount()`

```
nvgraphStatus_t
nvgraphTriangleCount(nvgraphHandle_t handle,
                    const nvgraphGraphDescr_t graph_descr,
                    uint64_t* result);
```

The triangles counting algorithm calculates number of unique triangles formed by graph edges. Algorithm works on the undirected graphs and graph structure in parameter should store only lower triangular of the adjacency matrix (no diagonal or self loops for vertices).

Input

| | |
|--------------------------|---|
| <code>handle</code> | nvGRAPH library handle. |
| <code>graph_descr</code> | nvGRAPH graph descriptor, should contain graph connectivity information in <code>NVGRAPH_CSR_32</code> or <code>NVGRAPH_CSC_32</code> format which should contain only lower triangular of adjacency matrix (no diagonal or self loops for vertices). |

Output

| | |
|---------------------|--|
| <code>result</code> | Number of triangles for provided graph will be stored in this parameter. |
|---------------------|--|

Return Values

| | |
|--|--|
| <code>NVGRAPH_STATUS_SUCCESS</code> | Success. |
| <code>NVGRAPH_STATUS_INVALID_VALUE</code> | Bad parameter(s), for example unsupported topology or topology is not stored in the graph. |
| <code>NVGRAPH_STATUS_MAPPING_ERROR</code> | Incorrect graph handle parameter. |
| <code>NVGRAPH_STATUS_INTERNAL_ERROR</code> | An internal operation failed. |

2.26. Function nvgraphStatusGetString()

```
const char*
nvgraphStatusGetString(nvgraphStatus_t status);
```

Gets string description for the nvGRAPH C API statuses.

Input

| | |
|---------------|---|
| status | Status returned from one of the C API functions |
|---------------|---|

Return Values

| |
|--|
| Pointer to the string with the text description of the C API status. |
|--|

2.27. nvGRAPH Spectral Clustering API

This part describes nvGRAPH Spectral Clustering API types, parameters and functions.

Graph clustering consists in grouping vertices based on their characteristics such that there is high intra-cluster similarity and low inter-cluster similarity. There are many ways of determining these groups. The spectral clustering scheme constructs a matrix, solves an associated eigenvalue problem, and extracts splitting information from the calculated eigenvectors.

nvGRAPH can compute cluster assignments with the spectral technique by using the modularity maximization or the minimum balanced cut formulation.

The modularity metric is based on the idea that similar vertices are connected by more edges in the current graph than if they were randomly connected. It measures how well a given clustering applies to a particular graph versus a random graph. The clustering is achieved by finding vertex assignments into clusters such that the modularity is maximized.

The balanced cut metric measures the size of each cluster and the volume of connections between clusters. It attempts to minimize the volume of connections between clusters relatively to their size.

It is important to outline the difference between the clustering and partitioning. Clustering is often used in analysis of graphs and networks, while partitioning is often used for load balancing across a fixed set of resources. For example, in partitioning the size of sub-graphs is specified and fixed while clustering focuses on finding tightly connected groups.

2.27.1. Structure SpectralClusteringParameter

```
struct SpectralClusteringParameter {
    int n_clusters;
    int n_eig_vects;
    nvgraphSpectralClusteringType_t algorithm;
    float evs_tolerance;
    int evs_max_iter;
    float kmean_tolerance;
    int kmean_max_iter;
};
```

Structure fields

| | |
|------------------------|---|
| n_clusters | Number of clusters. Should be at least 2. |
| n_eig_vects | Number of eigenvectors. Should be at least 2 and at most n_clusters . |
| algorithm | <p>nvGRAPH has pre-defined algorithms in nvgraphSpectralClusteringType_t which can take the following values:</p> <p>NVGRAPH_MODULARITY_MAXIMIZATION : maximize modularity with Lanczos solver.</p> <p>NVGRAPH_BALANCED_CUT_LANCZOS : minimize balanced cut with Lanczos solver.</p> <p>NVGRAPH_BALANCED_CUT_LOBPCG: minimize balanced cut with LOBPCG solver.</p> |
| evs_tolerance | Set the approximation tolerance for the eigensolver, this parameter should be a small magnitude value. The lower the tolerance the better the approximation. If this value is 0.0f , nvGRAPH will use the default value which is 1.0E-3 . Setting a low tolerance (less than 1.0E-4) may lead to divergence due to numerical roundoff. Usually values between 0.01 and 0.0001 are acceptable for spectral clustering. |
| evs_max_iter | The maximum number of eigensolver iterations before an answer is returned. This can be used to limit the execution time and do an early exit before the it reaches the convergence tolerance. If this value is lower or equal to 0 nvGRAPH will use the default value, which is 4000. |
| kmean_tolerance | Set the approximation tolerance for the k-means algorithm, this parameter should be a small magnitude value. The lower the tolerance the better the approximation. If this value is 0.0f , nvGRAPH will use the default value which is 1.0E-2 . Usually values between 0.01 and 0.001 are acceptable for spectral clustering. |
| kmean_max_iter | The maximum number of k-means iterations before an answer is returned. This can be used to limit the execution time and do an early exit before the it reaches the convergence tolerance. If this value is lower or equal to 0 nvGRAPH will use the default value, which is 200. |

2.27.2. Function nvgraphSpectralClustering()

```
nvgraphStatus_t
nvgraphSpectralClustering(nvgraphHandle_t handle,
                          const nvgraphGraphDescr_t descrG,
                          const size_t weight_index,
                          const struct SpectralClusteringParameter *params,
                          int* clustering,
                          void* eig_vals,
                          void* eig_vects);
```

Assign vertices to groups such as intra-group connections are strong and/or inter-groups connections are weak using spectral technique. This function supports weighted, undirected graphs.

Input

| | |
|---------------------|---|
| handle | nvGRAPH library handle. |
| descrG | nvGRAPH graph descriptor, should contain the connectivity information of an undirected graph in <code>NVGRAPH_CSR_32</code> , and at least 1 edge set. In <code>NVGRAPH_CSR_32</code> topology, an undirected edge between vertices u and v is represented as a directed edge in both direction (u,v) and (v,u) . |
| weight_index | Index of the edge set for the connection strength. For unweighted graphs, the edge set should be uniform and different than 0. |
| *params | Spectral clustering solver parameters stored in <code>struct SpectralClusteringParameter</code> |

Output

| | |
|--------------------|--|
| *clustering | For each vertex i , <code>clustering[i]</code> contains the cluster assignment computed by the spectral clustering algorithm. Possible values are between 0 and <code>n_clusters-1</code> . |
| *eig_vals | <p>If <code>algorithm=NVGRAPH_MODULARITY_MAXIMIZATION</code>, <code>*eig_vals</code> contains the <code>n_eig_vects</code> largest eigenvalues of the Modularity matrix.</p> <p>If <code>algorithm=NVGRAPH_BALANCED_CUT_LANCZOS</code> or <code>algorithm=NVGRAPH_BALANCED_CUT_LOBPCG</code>, <code>*eig_vals</code> contains the <code>n_eig_vects</code> smallest eigenvalues of the graph Laplacian matrix.</p> <p>The datatype is the same as the edge set at <code>weight_index</code></p> |
| *eig_vects | <p>If <code>algorithm=NVGRAPH_MODULARITY_MAXIMIZATION</code>, <code>*eig_vects</code> contains the eigeninformation based on the <code>n_eig_vects</code> largest eigenvectors of the Modularity matrix. The eigenvectors are stored in row-major order and normalized by row and by column.</p> <p>If <code>algorithm=NVGRAPH_BALANCED_CUT_LANCZOS</code> or <code>algorithm=NVGRAPH_BALANCED_CUT_LOBPCG</code>, <code>*eig_vects</code> contains the eigeninformation based on the <code>n_eig_vects</code> smallest eigenvectors of the graph Laplacian matrix.</p> |

| | |
|--|---|
| | The eigenvectors are stored in column-major order and normalized by column. |
| | The datatype is the same as the edge set at weight_index |

Return Values

| | |
|--|---|
| NVGRAPH_STATUS_SUCCESS | The spectral clustering algorithm found a solution. |
| NVGRAPH_STATUS_NOT_CONVERGED | The eigensolver of the k-means algorithm did not reach the desired tolerance after max_iter iterations. |
| NVGRAPH_STATUS_INVALID_VALUE | Bad parameter(s). |
| NVGRAPH_STATUS_TYPE_NOT_SUPPORTED | The type of at least one vertex or edge set is not supported. Currently we support float and double type values. |
| NVGRAPH_STATUS_GRAPH_TYPE_NOT_SUPPORTED | The type of the graph is not NVGRAPH_CSR_32 . |

2.27.3. Function nvgraphAnalyzeClustering()

```
nvgraphStatus_t
nvgraphAnalyzeClustering( nvgraphHandle_t handle,
                          const nvgraphGraphDescr_t descrG,
                          const size_t weight_index,
                          const int n_clusters,
                          const int* clustering,
                          nvgraphClusteringMetric_t metric,
                          float * score);
```

Measure the clustering quality according to a metric. This function supports weighted, undirected graphs.

Input

| | |
|---------------------|--|
| handle | nvGRAPH library handle. |
| descrG | nvGRAPH graph descriptor, should contain the connectivity information of an undirected graph in NVGRAPH_CSR_32 , and at least 1 edge set. In NVGRAPH_CSR_32 topology, an undirected edge between vertices <i>u</i> and <i>v</i> is represented as a directed edge in both direction (<i>u,v</i>) and (<i>v,u</i>). |
| weight_index | Index of the edge set for the connection strength. For unweighted graphs, the edge set should be uniform and different than 0. |
| n_clusters | Total number of clusters in *clustering . |
| *clustering | Cluster assignments to be measured. |
| metric | nvGRAPH has pre-defined metrics in nvgraphClusteringMetric_t which can take the following values: NVGRAPH_MODULARITY : modularity clustering score telling how good the clustering is compared to random assignments. NVGRAPH_EDGE_CUT : total number of edges between clusters. |

| | |
|--|---|
| | NVGRAPH_RATIO_CUT : sum for all clusters of the number of edges going outside of the cluster divided by the number of vertices inside the cluster. |
|--|---|

Output

| | |
|---------------|--|
| *score | Pointer to the score that measures the quality of the cluster assignments according to the metric . |
|---------------|--|

Return Values

| | |
|--|--|
| NVGRAPH_STATUS_SUCCESS | Success. |
| NVGRAPH_STATUS_INVALID_VALUE | Bad parameter(s). |
| NVGRAPH_STATUS_TYPE_NOT_SUPPORTED | The type of at least one vertex or edge set is not supported. Currently we support float and double type values. |
| NVGRAPH_STATUS_GRAPH_TYPE_NOT_SUPPORTED | The type of the graph is not NVGRAPH_CSR_32 . |

2.28. nvGRAPH Traversal API

This chapter describes nvGRAPH Traversal API types, parameters and functions.

2.28.1. nvGRAPH traversal algorithms

Enum which is used as a parameter to nvgraphTraversal API and which tells nvGRAPH which traversal algorithm to use.

```
typedef enum {
    NVGRAPH_TRAVERSAL_BFS=0
} nvgraphTraversal_t;
```

Algorithms

| | |
|------------------------------|---|
| NVGRAPH_TRAVERSAL_BFS | Use Breadth First Search algorithm for graph traversal. |
|------------------------------|---|

2.28.2. nvgraphTraversalParameter_t parameter struct

```
typedef struct {
    size_t pad[128];
} nvgraphTraversalParameter_t;
```

Array of bytes that is used for passing parameters to various algorithms of nvgraphTraversal API. Should be accessed through the corresponding Get/Set function.

2.28.3. nvgraphTraversalParameterInit()

```
nvgraphStatus_t nvgraphTraversalParameterInit(
    nvgraphTraversalParameter_t *param);
```

This function initializes **traversal parameter** struct with default values.

| parameter | Memory | In/out | Meaning |
|-----------|--------|--------|--|
| param | host | output | The pointer to the structure traversal parameters. |

Status Returned

| | |
|------------------------------|---|
| NVGRAPH_STATUS_SUCCESS | The structure was initialized successfully. |
| NVGRAPH_STATUS_INVALID_VALUE | Invalid parameter. |

2.28.4. nvgraphTraversalSetDistancesIndex()

```
nvgraphStatus_t
nvgraphTraversalSetDistancesIndex(
    nvgraphTraversalParameter_t *param,
    const size_t value);
```

This function sets index of vertex data where distances values will be stored. After successfull traversal this `vertex_data[vertex_id]` will be equal to shortest distance from traversal source_vertex to vertex_id. Distance to unreachable vertices is $2^{31}-1$ and for source vertex itself - 0.

| parameter | Memory | In/out | Meaning |
|-----------|--------|--------------|--|
| param | host | input/output | The pointer to the structure traversal parameters. |
| value | host | input | Index of the vertex data where to save distances. |

Status Returned

| | |
|------------------------------|--|
| CUSOLVER_STATUS_SUCCESS | Parameter was set successfully in the structure. |
| NVGRAPH_STATUS_INVALID_VALUE | Invalid parameter. |

2.28.5. nvgraphTraversalGetDistancesIndex()

```
nvgraphStatus_t
nvgraphTraversalGetDistancesIndex(
    nvgraphTraversalParameter_t *param,
    const size_t value);
```

This function retrieves index of vertex data for distances from the traversal parameters structure. See `nvgraphTraversalSetDistancesIndex()` description for predecessors data meaning.

| parameter | Memory | In/out | Meaning |
|-----------|--------|--------|---|
| param | host | input | The pointer to the structure traversal parameters. |
| value | host | output | Index of the vertex data for distances saved in the param struct. |

Status Returned

| | |
|------------------------------|---|
| CUSOLVER_STATUS_SUCCESS | Parameter was retrieved successfully from the params structure. |
| NVGRAPH_STATUS_INVALID_VALUE | Invalid parameter. |

2.28.6. nvgraphTraversalSetPredecessorsIndex()

```
nvgraphStatus_t
nvgraphTraversalSetPredecessorsIndex(
    nvgraphTraversalParameter_t *param,
    const size_t value);
```

This function sets index of vertex data where predecessor of the vertex in shortest path will be stored. After successful traversal if `vertex_data[vertex_id]=prev_vertex_id` that means that traversal path from **source vertex** to **vertex_id** contains edge from **prev_vertex_id** to **vertex_id**. Predecessor for source vertex and unreachable vertices is -1.

| parameter | Memory | In/out | Meaning |
|-----------|--------|--------------|---|
| param | host | input/output | The pointer to the structure traversal parameters. |
| value | host | input | Index of the vertex data where predecessors will be stored. |

Status Returned

| | |
|------------------------------|--|
| CUSOLVER_STATUS_SUCCESS | Parameter was set successfully in the structure. |
| NVGRAPH_STATUS_INVALID_VALUE | Invalid parameter. |

2.28.7. nvgraphTraversalGetPredecessorsIndex()

```
nvgraphStatus_t
nvgraphTraversalGetPredecessorsIndex(
    nvgraphTraversalParameter_t *param,
    const size_t value);
```

This function retrieves index of vertex data for predecessors from traversal parameters structure. See **nvgraphTraversalSetPredecessorsIndex()** description for predecessors data meaning.

| parameter | Memory | In/out | Meaning |
|-----------|--------|--------|--|
| param | host | input | The pointer to the structure traversal parameters. |
| value | host | output | Index of the vertex data for predecessors saved in the param struct. |

Status Returned

| | |
|------------------------------|---|
| CUSOLVER_STATUS_SUCCESS | Parameter was retrieved successfully from the params structure. |
| NVGRAPH_STATUS_INVALID_VALUE | Invalid parameter. |

2.28.8. nvgraphTraversalSetEdgeMaskIndex()

```
nvgraphStatus_t
nvgraphTraversalSetEdgeMaskIndex(
    nvgraphTraversalParameter_t *param,
    const size_t value);
```

This function sets index of edge data where mask for edges will be stored. If for this set of data **edge_data[edge_id]** is 0 then traversal path will not advance using this edge.

| parameter | Memory | In/out | Meaning |
|-----------|--------|--------------|--|
| param | host | input/output | The pointer to the structure traversal parameters. |
| value | host | input | Index of the edge data where edge mask will be stored. |

Status Returned

| | |
|------------------------------|--|
| CUSOLVER_STATUS_SUCCESS | Parameter was set successfully in the structure. |
| NVGRAPH_STATUS_INVALID_VALUE | Invalid parameter. |

2.28.9. nvgraphTraversalGetEdgeMaskIndex()

```
nvgraphStatus_t
nvgraphTraversalGetEdgeMaskIndex(
    nvgraphTraversalParameter_t *param,
    const size_t value);
```

This function retrieves index of an edge data for edge mask from traversal parameters structure. If for this set of data **edge_data[edge_id]** is 0 then traversal path will not advance using this edge.

| parameter | Memory | In/out | Meaning |
|-----------|--------|--------|---|
| param | host | input | The pointer to the structure traversal parameters. |
| value | host | output | Index of the edge data for edge mask saved in the param struct. |

Status Returned

| | |
|------------------------------|---|
| CUSOLVER_STATUS_SUCCESS | Parameter was retrieved successfully from the params structure. |
| NVGRAPH_STATUS_INVALID_VALUE | Invalid parameter. |

2.28.10. nvgraphTraversalSetUndirectedFlag()

```
nvgraphStatus_t
nvgraphTraversalSetUndirectedFlag(
    nvgraphTraversalParameter_t *param,
    const size_t value);
```

This function sets flag specifying whether processing will be performed on the undirected graph or not.

| parameter | Memory | In/out | Meaning |
|-----------|--------|--------------|--|
| param | host | input/output | The pointer to the structure traversal parameters. |
| value | host | input | 1 if graph is undirected, 0 - otherwise. |

Status Returned

| | |
|------------------------------|--|
| CUSOLVER_STATUS_SUCCESS | Parameter was set successfully in the structure. |
| NVGRAPH_STATUS_INVALID_VALUE | Invalid parameter. |

2.28.11. nvgraphTraversalGetUndirectedFlag()

```
nvgraphStatus_t
nvgraphTraversalGetUndirectedFlag(
    nvgraphTraversalParameter_t *param,
    const size_t value);
```

This function retrieves graph direction property flag from traversal parameters structure.

| parameter | Memory | In/out | Meaning |
|-----------|--------|--------|--|
| param | host | input | The pointer to the structure traversal parameters. |
| value | host | output | 1 if graph is undirected, 0 - otherwise. |

Status Returned

| | |
|------------------------------|---|
| CUSOLVER_STATUS_SUCCESS | Parameter was retrieved successfully from the params structure. |
| NVGRAPH_STATUS_INVALID_VALUE | Invalid parameter. |

2.28.12. nvgraphTraversalSetAlpha()

```
nvgraphStatus_t
nvgraphTraversalSetAlpha(
    nvgraphTraversalParameter_t *param,
    const size_t value);
```

This function sets **alpha** parameter for BFS. See <http://parlab.eecs.berkeley.edu/sites/all/parlab/files/main.pdf> paper for the details on meaning of this parameter.

| parameter | Memory | In/out | Meaning |
|-----------|--------|--------------|--|
| param | host | input/output | The pointer to the structure traversal parameters. |
| value | host | input | Alpha parameter to be set. |

Status Returned

| | |
|------------------------------|--|
| CUSOLVER_STATUS_SUCCESS | Parameter was set successfully in the structure. |
| NVGRAPH_STATUS_INVALID_VALUE | Invalid parameter. |

2.28.13. nvgraphTraversalGetAlpha()

```
nvgraphStatus_t
nvgraphTraversalGetAlpha(
    nvgraphTraversalParameter_t *param,
    const size_t value);
```

This function retrieves **alpha** BFS parameter from traversal parameters structure. See <http://parlab.eecs.berkeley.edu/sites/all/parlab/files/main.pdf> paper for the details on meaning of this parameter.

| parameter | Memory | In/out | Meaning |
|-----------|--------|--------|--|
| param | host | input | The pointer to the structure traversal parameters. |
| value | host | output | Alpha parameter saved in the traversal parameters structure. |

Status Returned

| | |
|------------------------------|---|
| CUSOLVER_STATUS_SUCCESS | Parameter was retrieved successfully from the params structure. |
| NVGRAPH_STATUS_INVALID_VALUE | Invalid parameter. |

2.28.14. nvgraphTraversalSetBeta()

```
nvgraphStatus_t
nvgraphTraversalSetBeta(
    nvgraphTraversalParameter_t *param,
    const size_t value);
```

This function sets **beta** parameter for BFS traversal. See <http://parlab.eecs.berkeley.edu/sites/all/parlab/files/main.pdf> paper for the details on meaning of this parameter.

| parameter | Memory | In/out | Meaning |
|-----------|--------|--------------|--|
| param | host | input/output | The pointer to the structure traversal parameters. |
| value | host | input | Beta parameter to be set. |

Status Returned

| | |
|------------------------------|--|
| CUSOLVER_STATUS_SUCCESS | Parameter was set successfully in the structure. |
| NVGRAPH_STATUS_INVALID_VALUE | Invalid parameter. |

2.28.15. nvgraphTraversalGetBeta()

```
nvgraphStatus_t
nvgraphTraversalGetBeta(
    nvgraphTraversalParameter_t *param,
    const size_t value);
```

This function retrieves **beta** BFS parameter from traversal parameters structure. See <http://parlab.eecs.berkeley.edu/sites/all/parlab/files/main.pdf> paper for the details on meaning of this parameter.

| parameter | Memory | In/out | Meaning |
|-----------|--------|--------|--|
| param | host | input | The pointer to the structure traversal parameters. |
| value | host | output | Beta parameter saved in the traversal parameters structure. |

Status Returned

| | |
|------------------------------|---|
| CUSOLVER_STATUS_SUCCESS | Parameter was retrieved successfully from the params structure. |
| NVGRAPH_STATUS_INVALID_VALUE | Invalid parameter. |

2.28.16. nvgraphTraversal()

```
nvgraphStatus_t
nvgraphTraversal(nvgraphHandle_t handle,
    const nvgraphGraphDescr_t descrG,
    const nvgraphTraversal_t traversalT,
    const int *source_vert,
    const nvgraphTraversalParameter_t params);
```

This functions performs traversal of the graph specified by **descrG** descriptor from the **source_vert** using traversal **params** with algorithm provided in **traversalT**. See corresponding helper function description to get details on **params** structure meaning for each of the algorithm.

| parameter | Memory | In/out | Meaning |
|-------------|--------|--------------|---|
| descrG | host | input/output | Graph descriptor handle on which traversal will be performed. |
| traversalT | host | input | Traversal algorithm to be used. |
| source_vert | host | input | Pointer to the starting vertex for the traversal. |
| params | host | input | Traversal parameters structure. Use other helper functions to modify parameters values. |

Status Returned

| | |
|-------------------------|---------------------------------------|
| CUSOLVER_STATUS_SUCCESS | Traversal was performed successfully. |
|-------------------------|---------------------------------------|

| | |
|-------------------------------|-------------------------------|
| NVGRAPH_STATUS_INVALID_VALUE | Invalid parameter. |
| NVGRAPH_STATUS_INTERNAL_ERROR | An internal operation failed. |

Chapter 3.

NVGRAPH CODE EXAMPLES

This chapter provides simple examples.

3.1. nvGRAPH convert topology example

```

void check(nvgraphStatus_t status) {
    if (status != NVGRAPH_STATUS_SUCCESS) {
        printf("ERROR : %d\n", status);
        exit(0);
    }
}

int main(int argc, char **argv) {
    size_t n = 6, nnz = 10;
    // nvgraph variables
    nvgraphHandle_t handle;
    nvgraphCSCTopology32I_t CSC_input;
    nvgraphCSRTopology32I_t CSR_output;
    float *src_weights_d, *dst_weights_d;
    cudaDataType_t edge_dimT = CUDA_R_32F;
    // Allocate source data
    CSC_input = (nvgraphCSCTopology32I_t) malloc(sizeof(struct
nvgraphCSCTopology32I_st));
    CSC_input->nvertices = n; CSC_input->nedges = nnz;
    cudaMalloc( (void**)&(CSC_input->destination_offsets), (n+1)*sizeof(int));
    cudaMalloc( (void**)&(CSC_input->source_indices), nnz*sizeof(int));
    cudaMalloc( (void**)&src_weights_d, nnz*sizeof(float));
    // Copy source data
    float src_weights_h[] = {0.333333f, 0.5f, 0.333333f, 0.5f, 0.5f, 1.0f,
0.333333f, 0.5f, 0.5f, 0.5f};
    int destination_offsets_h[] = {0, 1, 3, 4, 6, 8, 10};
    int source_indices_h[] = {2, 0, 2, 0, 4, 5, 2, 3, 3, 4};
    cudaMemcpy(CSC_input->destination_offsets, destination_offsets_h, (n
+1)*sizeof(int), cudaMemcpyDefault);
    cudaMemcpy(CSC_input->source_indices, source_indices_h, nnz*sizeof(int),
cudaMemcpyDefault);
    cudaMemcpy(src_weights_d, src_weights_h, nnz*sizeof(float),
cudaMemcpyDefault);
    // Allocate destination data
    CSR_output = (nvgraphCSRTopology32I_t) malloc(sizeof(struct
nvgraphCSRTopology32I_st));
    cudaMalloc( (void**)&(CSR_output->source_offsets), (n+1)*sizeof(int));
    cudaMalloc( (void**)&(CSR_output->destination_indices), nnz*sizeof(int));
    cudaMalloc( (void**)&dst_weights_d, nnz*sizeof(float));
    // Starting nvgraph and convert
    check(nvgraphCreate (&handle));
    check(nvgraphConvertTopology(handle, NVGRAPH_CSC_32, CSC_input,
src_weights_d,
        &edge_dimT, NVGRAPH_CSR_32, CSR_output, dst_weights_d));
    // Free memory
    check(nvgraphDestroy(handle));
    cudaFree(CSC_input->destination_offsets);
    cudaFree(CSC_input->source_indices);
    cudaFree(CSR_output->source_offsets);
    cudaFree(CSR_output->destination_indices);
    cudaFree(src_weights_d);
    cudaFree(dst_weights_d);
    free(CSC_input);
    free(CSR_output);
    return 0;
}

```

3.2. nvGRAPH convert graph example

```

void check(nvgraphStatus_t status) {
    if (status != NVGRAPH_STATUS_SUCCESS) {
        printf("ERROR : %d\n", status);
        exit(0);
    }
}

int main(int argc, char **argv) {
    size_t n = 6, nnz = 10, vert_sets = 2, edge_sets = 1;
    // nvgraph variables
    nvgraphHandle_t handle; nvgraphGraphDescr_t src_csc_graph;
    nvgraphCSCTopology32I_t CSC_input;
    cudaDataType_t edge_dimT = CUDA_R_32F;
    cudaDataType_t* vertex_dimT;
    // Allocate host data
    float *pr_1 = (float*)malloc(n*sizeof(float));
    void **vertex_dim = (void**)malloc(vert_sets*sizeof(void*));
    vertex_dimT = (cudaDataType_t*)malloc(vert_sets*sizeof(cudaDataType_t));
    CSC_input = (nvgraphCSCTopology32I_t) malloc(sizeof(struct
nvgraphCSCTopology32I_st));
    // Initialize host data
    float weights_h[] = {0.333333f, 0.5f, 0.333333f, 0.5f, 0.5f, 1.0f,
0.333333f, 0.5f, 0.5f, 0.5f};
    int destination_offsets_h[] = {0, 1, 3, 4, 6, 8, 10};
    int source_indices_h[] = {2, 0, 2, 0, 4, 5, 2, 3, 3, 4};
    float bookmark_h[] = {0.0f, 1.0f, 0.0f, 0.0f, 0.0f, 0.0f};
    vertex_dim[0] = (void*)bookmark_h; vertex_dim[1] = (void*)pr_1;
    vertex_dimT[0] = CUDA_R_32F; vertex_dimT[1] = CUDA_R_32F, vertex_dimT[2] =
CUDA_R_32F;
    // Starting nvgraph
    check(nvgraphCreate (&handle));
    check(nvgraphCreateGraphDescr (handle, &src_csc_graph));
    CSC_input->nvertices = n; CSC_input->nedges = nnz;
    CSC_input->destination_offsets = destination_offsets_h;
    CSC_input->source_indices = source_indices_h;
    // Set graph connectivity and properties (transfers)
    check(nvgraphSetGraphStructure(handle, src_csc_graph, (void*)CSC_input,
NVGRAPH_CSC_32));
    check(nvgraphAllocateVertexData(handle, src_csc_graph, vert_sets,
vertex_dimT));
    check(nvgraphAllocateEdgeData (handle, src_csc_graph, edge_sets,
&edge_dimT));
    for (int i = 0; i < 2; ++i)
        check(nvgraphSetVertexData(handle, src_csc_graph, vertex_dim[i], i));
    check(nvgraphSetEdgeData(handle, src_csc_graph, (void*)weights_h, 0));
    // Convert to CSR graph
    nvgraphGraphDescr_t dst_csr_graph;
    check(nvgraphCreateGraphDescr (handle, &dst_csr_graph));
    check(nvgraphConvertGraph(handle, src_csc_graph, dst_csr_graph,
NVGRAPH_CSR_32));
    check(nvgraphDestroyGraphDescr(handle, src_csc_graph));
    check(nvgraphDestroyGraphDescr(handle, dst_csr_graph));
    check(nvgraphDestroy(handle));
    free(pr_1); free(vertex_dim); free(vertex_dimT);
    free(CSC_input);
    return 0;
}

```

3.3. nvGRAPH pagerank example

```

void check(nvgraphStatus_t status) {
    if (status != NVGRAPH_STATUS_SUCCESS) {
        printf("ERROR : %d\n", status);
        exit(0);
    }
}

int main(int argc, char **argv) {
    size_t n = 6, nnz = 10, vert_sets = 2, edge_sets = 1;
    float alpha1 = 0.9f; void *alpha1_p = (void *) &alpha1;
    // nvgraph variables
    nvgraphHandle_t handle; nvgraphGraphDescr_t graph;
    nvgraphCSCTopology32I_t CSC_input;
    cudaDataType_t edge_dimT = CUDA_R_32F;
    cudaDataType_t *vertex_dimT;
    // Allocate host data
    float *pr_1 = (float *) malloc(n * sizeof(float));
    void **vertex_dim = (void **) malloc(vert_sets * sizeof(void *));
    vertex_dimT = (cudaDataType_t *) malloc(vert_sets * sizeof(cudaDataType_t));
    CSC_input = (nvgraphCSCTopology32I_t) malloc(sizeof(struct
nvgraphCSCTopology32I_st));
    // Initialize host data
    float weights_h[] = {0.333333f, 0.5f, 0.333333f, 0.5f, 0.5f, 1.0f,
0.333333f, 0.5f, 0.5f, 0.5f};
    int destination_offsets_h[] = {0, 1, 3, 4, 6, 8, 10};
    int source_indices_h[] = {2, 0, 2, 0, 4, 5, 2, 3, 3, 4};
    float bookmark_h[] = {0.0f, 1.0f, 0.0f, 0.0f, 0.0f, 0.0f};
    vertex_dim[0] = (void *) bookmark_h; vertex_dim[1] = (void *) pr_1;
    vertex_dimT[0] = CUDA_R_32F; vertex_dimT[1] = CUDA_R_32F, vertex_dimT[2] =
CUDA_R_32F;
    // Starting nvgraph
    check(nvgraphCreate (&handle));
    check(nvgraphCreateGraphDescr (handle, &graph));
    CSC_input->nvertices = n; CSC_input->nedges = nnz;
    CSC_input->destination_offsets = destination_offsets_h;
    CSC_input->source_indices = source_indices_h;
    // Set graph connectivity and properties (transfers)
    check(nvgraphSetGraphStructure(handle, graph, (void *) CSC_input,
NVGRAPH_CSC_32));
    check(nvgraphAllocateVertexData(handle, graph, vert_sets, vertex_dimT));
    check(nvgraphAllocateEdgeData (handle, graph, edge_sets, &edge_dimT));
    for (int i = 0; i < 2; ++i)
        check(nvgraphSetVertexData(handle, graph, vertex_dim[i], i));
    check(nvgraphSetEdgeData(handle, graph, (void *) weights_h, 0));

    check(nvgraphPagerank(handle, graph, 0, alpha1_p, 0, 0, 1, 0.0f, 0));
    // Get result
    check(nvgraphGetVertexData(handle, graph, vertex_dim[1], 1));
    check(nvgraphDestroyGraphDescr(handle, graph));
    check(nvgraphDestroy(handle));
    free(pr_1); free(vertex_dim); free(vertex_dimT);
    free(CSC_input);
    return 0;
}

```

3.4. nvGRAPH SSSP example

```

void check(nvgraphStatus_t status) {
    if (status != NVGRAPH_STATUS_SUCCESS) {
        printf("ERROR : %d\n", status);
        exit(0);
    }
}

int main(int argc, char **argv) {
    const size_t n = 6, nnz = 10, vertex_numsets = 1, edge_numsets = 1;
    float *sssp_1_h;
    void** vertex_dim;
    // nvgraph variables
    nvgraphStatus_t status; nvgraphHandle_t handle;
    nvgraphGraphDescr_t graph;
    nvgraphCSCTopology32I_t CSC_input;
    cudaDataType_t edge_dimT = CUDA_R_32F;
    cudaDataType_t* vertex_dimT;
    // Init host data
    sssp_1_h = (float*)malloc(n*sizeof(float));
    vertex_dim = (void**)malloc(vertex_numsets*sizeof(void*));
    vertex_dimT =
    (cudaDataType_t*)malloc(vertex_numsets*sizeof(cudaDataType_t));
    CSC_input = (nvgraphCSCTopology32I_t) malloc(sizeof(struct
    nvgraphCSCTopology32I_st));
    vertex_dim[0] = (void*)sssp_1_h; vertex_dimT[0] = CUDA_R_32F;
    float weights_h[] = {0.333333, 0.5, 0.333333, 0.5, 0.5, 1.0, 0.333333, 0.5,
    0.5, 0.5};
    int destination_offsets_h[] = {0, 1, 3, 4, 6, 8, 10};
    int source_indices_h[] = {2, 0, 2, 0, 4, 5, 2, 3, 3, 4};
    check(nvgraphCreate(&handle));
    check(nvgraphCreateGraphDescr (handle, &graph));
    CSC_input->nvertices = n; CSC_input->nedges = nnz;
    CSC_input->destination_offsets = destination_offsets_h;
    CSC_input->source_indices = source_indices_h;
    // Set graph connectivity and properties (transfers)
    check(nvgraphSetGraphStructure (handle, graph, (void*)CSC_input,
    NVGRAPH_CSC_32));
    check(nvgraphAllocateVertexData (handle, graph, vertex_numsets,
    vertex_dimT));
    check(nvgraphAllocateEdgeData (handle, graph, edge_numsets, &edge_dimT));
    check(nvgraphSetEdgeData (handle, graph, (void*)weights_h, 0));
    // Solve
    int source_vert = 0;
    check(nvgraphSssp(handle, graph, 0, &source_vert, 0));
    // Get and print result
    check(nvgraphGetVertexData (handle, graph, (void*)sssp_1_h, 0));
    //Clean
    free(sssp_1_h); free(vertex_dim);
    free(vertex_dimT); free(CSC_input);
    check(nvgraphDestroyGraphDescr (handle, graph));
    check(nvgraphDestroy (handle));
    return 0;
}

```

3.5. nvGRAPH Semi-Ring SPMV example

```

void check(nvgraphStatus_t status) {
    if (status != NVGRAPH_STATUS_SUCCESS) {
        printf("ERROR : %d\n", status);
        exit(0);
    }
}

int main(int argc, char **argv) {
    size_t n = 5, nnz = 10, vertex_numsets = 2, edge_numsets = 1;
    float alpha = 1.0, beta = 0.0;
    void *alpha_p = (void *)&alpha, *beta_p = (void *)&beta;
    void** vertex_dim;
    cudaDataType_t edge_dimT = CUDA_R_32F;
    cudaDataType_t* vertex_dimT;
    // nvgraph variables
    nvgraphStatus_t status; nvgraphHandle_t handle;
    nvgraphGraphDescr_t graph;
    nvgraphCSRTopology32I_t CSR_input;
    // Init host data
    vertex_dim = (void**)malloc(vertex_numsets*sizeof(void*));
    vertex_dimT =
    (cudaDataType_t*)malloc(vertex_numsets*sizeof(cudaDataType_t));
    CSR_input = (nvgraphCSRTopology32I_t) malloc(sizeof(struct
    nvgraphCSRTopology32I_st));
    float x_h[] = {1.1f, 2.2f, 3.3f, 4.4f, 5.5f};
    float y_h[] = {0.0f, 0.0f, 0.0f, 0.0f, 0.0f};
    vertex_dim[0] = (void*)x_h; vertex_dim[1] = (void*)y_h;
    vertex_dimT[0] = CUDA_R_32F; vertex_dimT[1] = CUDA_R_32F;
    float weights_h[] = {1.0f, 4.0f, 2.0f, 3.0f, 5.0f, 7.0f, 8.0f, 9.0f, 6.0f,
    1.5f};
    int source_offsets_h[] = {0, 2, 4, 7, 9, 10};
    int destination_indices_h[] = {0, 1, 1, 2, 0, 3, 4, 2, 4, 2};
    check(nvgraphCreate(&handle));
    check(nvgraphCreateGraphDescr(handle, &graph));
    CSR_input->nvertices = n; CSR_input->nedges = nnz;
    CSR_input->source_offsets = source_offsets_h;
    CSR_input->destination_indices = destination_indices_h;
    // Set graph connectivity and properties (transfers)
    check(nvgraphSetGraphStructure(handle, graph, (void*)CSR_input,
    NVGRAPH_CSR_32));
    check(nvgraphAllocateVertexData(handle, graph, vertex_numsets,
    vertex_dimT));
    for (int i = 0; i < vertex_numsets; ++i)
        check(nvgraphSetVertexData(handle, graph, vertex_dim[i], i));
    check(nvgraphAllocateEdgeData (handle, graph, edge_numsets, &edge_dimT));
    check(nvgraphSetEdgeData(handle, graph, (void*)weights_h, 0));
    // Solve
    check(nvgraphSrSpmv(handle, graph, 0, alpha_p, 0, beta_p, 1,
    NVGRAPH_PLUS_TIMES_SR));
    //Get result
    check(nvgraphGetVertexData(handle, graph, (void*)y_h, 1));
    //Clean
    check(nvgraphDestroyGraphDescr(handle, graph));
    check(nvgraphDestroy(handle));
    free(vertex_dim); free(vertex_dimT); free(CSR_input);
    return 0;
}

```

3.6. nvGRAPH Triangles Counting example

```
#include "stdlib.h"
#include "inttypes.h"
#include "stdio.h"

#include "nvgraph.h"

#define check( a ) \
{ \
    nvgraphStatus_t status = (a);\
    if ( (status) != NVGRAPH_STATUS_SUCCESS) {\
        printf("ERROR : %d in %s : %d\n", status, __FILE__ , __LINE__ );\
        exit(0);\
    } \
}

int main(int argc, char **argv)
{
    // nvgraph variables
    nvgraphHandle_t handle;
    nvgraphGraphDescr_t graph;
    nvgraphCSRTopology32I_t CSR_input;

    // Init host data
    CSR_input = (nvgraphCSRTopology32I_t) malloc(sizeof(struct
    nvgraphCSRTopology32I_st));

    // Undirected graph:
    // 0          2-----4
    //  \         / \     / \
    //   \       /   \   /   \
    //    \     /     \ /     \
    //     \   /       \ /       \
    //      1-----3-----5
    // 3 triangles
    // CSR of lower triangular of adjacency matrix:
    const size_t n = 6, nnz = 8;
    int source_offsets[] = {0, 0, 1, 2, 4, 6, 8};
    int destination_indices[] = {0, 1, 1, 2, 2, 3, 3, 4};

    check(nvgraphCreate(&handle));
    check(nvgraphCreateGraphDescr (handle, &graph));
    CSR_input->nvertices = n;
    CSR_input->nedges = nnz;
    CSR_input->source_offsets = source_offsets;
    CSR_input->destination_indices = destination_indices;
    // Set graph connectivity
    check(nvgraphSetGraphStructure(handle, graph, (void*)CSR_input,
    NVGRAPH_CSR_32));

    uint64_t trcount = 0;
    check(nvgraphTriangleCount(handle, graph, &trcount));
    printf("Triangles count: %" PRIu64 "\n", trcount);

    free(CSR_input);
    check(nvgraphDestroyGraphDescr(handle, graph));
    check(nvgraphDestroy(handle));
    return 0;
}
```

3.7. nvGRAPH Traversal example

```

void check_status(nvgraphStatus_t status){
    if ((int)status != 0) {
        printf("ERROR : %d\n",status);
        exit(0);
    }
}

int main(int argc, char **argv){
    //Example of graph (CSR format)
    const size_t n = 7, nnz = 12, vertex_numsets = 2, edge_numset = 0;
    int source_offsets_h[] = {0, 1, 3, 4, 6, 8, 10, 12};
    int destination_indices_h[] = {5, 0, 2, 0, 4, 5, 2, 3, 3, 4, 1, 5};
    //where to store results (distances from source) and where to store results
    (predecessors in search tree)
    int bfs_distances_h[n], bfs_predecessors_h[n];
    // nvgraph variables
    nvgraphStatus_t status;
    nvgraphHandle_t handle;
    nvgraphGraphDescr_t graph;
    nvgraphCSRTopology32I_t CSR_input;
    cudaDataType_t* vertex_dimT;
    size_t distances_index = 0;
    size_t predecessors_index = 1;
    vertex_dimT =
(cudaDataType_t*)malloc(vertex_numsets*sizeof(cudaDataType_t));
    vertex_dimT[distances_index] = CUDA_R_32I;
    vertex_dimT[predecessors_index] = CUDA_R_32I;
    //Creating nvgraph objects
    check_status(nvgraphCreate (&handle));
    check_status(nvgraphCreateGraphDescr (handle, &graph));
    // Set graph connectivity and properties (transfers)
    CSR_input = (nvgraphCSRTopology32I_t) malloc(sizeof(struct
nvgraphCSCTopology32I_st));
    CSR_input->nvertices = n;
    CSR_input->nedges = nnz;
    CSR_input->source_offsets = source_offsets_h;
    CSR_input->destination_indices = destination_indices_h;
    check_status(nvgraphSetGraphStructure(handle, graph, (void*)CSR_input,
NVGRAPH_CSR_32));
    check_status(nvgraphAllocateVertexData(handle, graph, vertex_numsets,
vertex_dimT));
    int source_vert = 1;
    //Setting the traversal parameters
    nvgraphTraversalParameter_t traversal_param;
    nvgraphTraversalParameterInit(&traversal_param);
    nvgraphTraversalSetDistancesIndex(&traversal_param, distances_index);
    nvgraphTraversalSetPredecessorsIndex(&traversal_param, predecessors_index);
    nvgraphTraversalSetUndirectedFlag(&traversal_param, false);
    //Computing traversal using BFS algorithm
    check_status(nvgraphTraversal(handle, graph, NVGRAPH_TRAVERSAL_BFS,
&source_vert, traversal_param));
    // Get result
    check_status(nvgraphGetVertexData(handle, graph, (void*)bfs_distances_h,
distances_index));
    check_status(nvgraphGetVertexData(handle, graph, (void*)bfs_predecessors_h,
predecessors_index));
    // expect bfs distances_h = {1 0 1 3 3 2 2147483647}
    for (int i = 0; i<n; i++) printf("Distance to vertex %d: %i\n",i,
bfs_distances_h[i]); printf("\n");
    // expect bfs predecessors = {1 -1 1 5 5 0 -1}
    for (int i = 0; i<n; i++) printf("Predecessor of vertex %d: %i\n",i,
bfs_predecessors_h[i]); printf("\n");
    free(vertex_dimT);
    free(CSR_input);
    check_status(nvgraphDestroyGraphDescr (handle, graph));
    check_status(nvgraphDestroy (handle));
    return 0;
}

```

Notice

ALL NVIDIA DESIGN SPECIFICATIONS, REFERENCE BOARDS, FILES, DRAWINGS, DIAGNOSTICS, LISTS, AND OTHER DOCUMENTS (TOGETHER AND SEPARATELY, "MATERIALS") ARE BEING PROVIDED "AS IS." NVIDIA MAKES NO WARRANTIES, EXPRESSED, IMPLIED, STATUTORY, OR OTHERWISE WITH RESPECT TO THE MATERIALS, AND EXPRESSLY DISCLAIMS ALL IMPLIED WARRANTIES OF NONINFRINGEMENT, MERCHANTABILITY, AND FITNESS FOR A PARTICULAR PURPOSE.

Information furnished is believed to be accurate and reliable. However, NVIDIA Corporation assumes no responsibility for the consequences of use of such information or for any infringement of patents or other rights of third parties that may result from its use. No license is granted by implication of otherwise under any patent rights of NVIDIA Corporation. Specifications mentioned in this publication are subject to change without notice. This publication supersedes and replaces all other information previously supplied. NVIDIA Corporation products are not authorized as critical components in life support devices or systems without express written approval of NVIDIA Corporation.

Trademarks

NVIDIA and the NVIDIA logo are trademarks or registered trademarks of NVIDIA Corporation in the U.S. and other countries. Other company and product names may be trademarks of the respective companies with which they are associated.

Copyright

© 2016-2019 NVIDIA Corporation. All rights reserved.